

STAT 542 Midterm Project Presentation

Group 8:

Group Member: Yefan Li, Jue Li, Yijin Wang, Xiao Ma, Wenxuan Gu, Pengru Lyu, Ziyi Xue, Yanqing Li

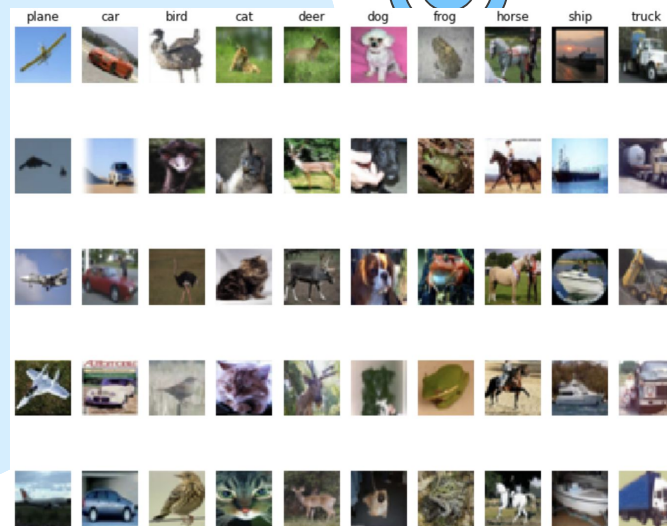


Dataset

...

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

The dataset is divided into five training batches and one test batch, each with 10,000 images. The test batch contains exactly 10,000 randomly-selected images from each class.



Support Vector Machines (SVM)

- SVM is a set of supervised learning methods that find the hyperplanes to separate datasets.
- Pros & Cons of SVM (with this dataset):

Pros: effective in high dimensional spaces, by mapping input data with different kernels

Cons:


- a. computationally intensive, especially we choose the k-fold CV method for turning parameters
- b. highly affected by the choice of the kernel

Preprocessing Data

Shape of the dataset:

- Change $32 * 32 * 3$ color image data into 1 Dimension
(potential problems caused by the really high dimensions)
- Scale the pixel value in $[0,1]$

Balance of the data:



SVM Model Parameters & Selection

Our SVM model is based on tuning the kernels and the parameters c and γ :

```
finalModel = SVC(C=1, gamma='scale', kernel='poly')
```

- Parameters:

Regularization parameter (c):

The strength of the regularization is inversely proportional to

C . Must be strictly positive. The penalty is a squared L_2 penalty.

γ : determine the width of the kernel (useless for linear kernel)


Kernel: 'linear'; 'poly'; 'rbf'; 'sigmoid'

Degree: Degree of the polynomial kernel function ('poly') and ignore by




all other kernel.

- Method:

- 1) Fixed γ value, different C value (1, 5, 10, 50, 100)
- 2) Fixed C value, different γ value ($1e-2$, $1e-3$, $1e-4$)



Results & Performance of SVM



```
linear_model = SVC(kernel='linear')
```

Linear SVM with accuracy around 30.31 %

10000 trained and test data

```
SVC(C=5, gamma='scale', kernel='rbf', degree=2)
```

Nonlinear (rbf) with parameter tuning:
Accuracy: 48.29%
Runtime: 1283.28 s (~21.4 min)

```
SVC(kernel='poly')
```

Nonlinear (poly) without parameter tuning:
Accuracy: 52.86%
Runtime: 9747.61 s (~3hr)

50000 trained data and 5000 test data



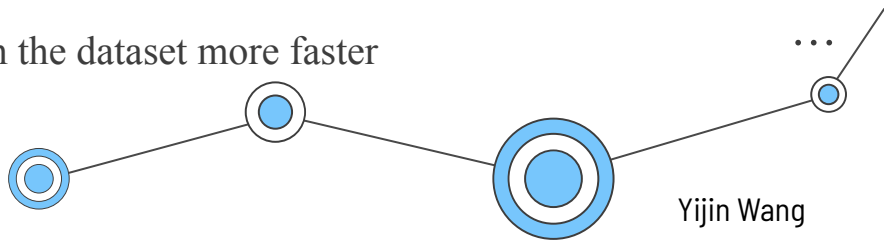
Difficulty & Future Work...




Difficulties & Problems:

- **Computation complexity:** for high dimensions ($p = 3072$) and large sample size (for both training and testing data)
- **Noise of dataset:** the images have much noises, which lead to the bad performance of SVM
- **Long runtime:** somewhat hard to run the whole 50,000 dataset at once therefore not having the most accurate result

Possible Solutions:

- **Computation complexity:** parallel processing
 - **Noise of dataset:** data smoothing methods, such as penalized models and resample with different ratios.
 - **Long runtime:** try to use CPU or GPU to run the dataset more faster
- 

Convolutional Neural Network(CNN)




What is neural network?




What is convolutional neural network (CNN)?



Why CNN?




Base Model Structure



Our Neural Network Model process can be explained as above

32 - 32 - 64 - 64



Activation Function Selection: Sigmoid VS ReLU





Figure: Sigmoid vs. ReLU

Source: https://miro.medium.com/v2/resize:fit:720/format:webp/1*XxxiA0jJvPrHEJHD4z893g.png

Sigmoid vs. ReLU

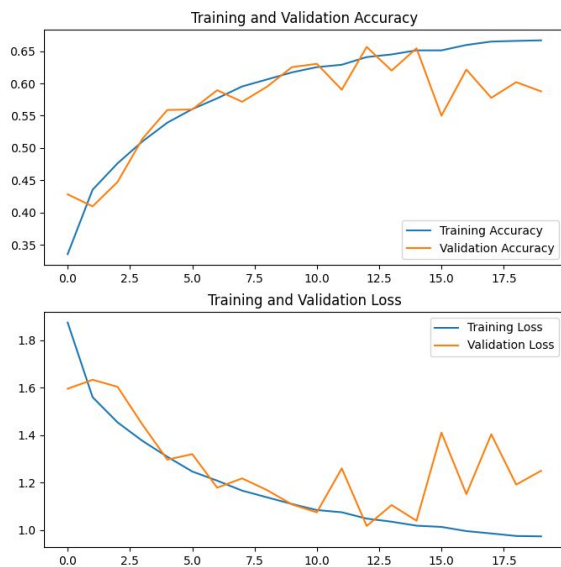


Figure: Sigmoid
Sigmoid Accuracy: 0.5877 (approx.)

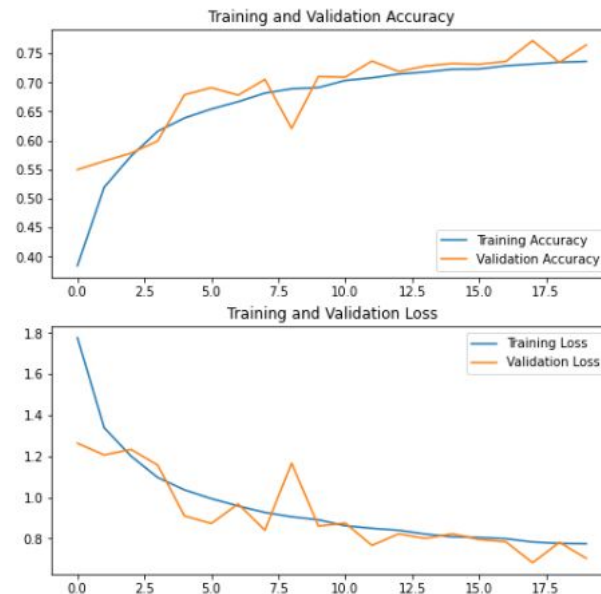


Figure: ReLU, accuracy: 0.7533 (approx.)



First Attempt: Add More Convolutional Layers



32-64-128

```
[4] amount_layers = len(model.layers)
amount_layers
```

17

```
model_summa = model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 64)	262208
batch_normalization_4 (Batch Normalization)	(None, 64)	256
dropout_2 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 10)	650

Total params: 329,450
 Trainable params: 328,938
 Non-trainable params: 512

```
[9] ayer_amount = len(model.layers)
ayer_amount
```

23

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_14 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_13 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_15 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_8 (Dropout)	(None, 16, 16, 32)	0
conv2d_14 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_16 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_15 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_17 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_9 (Dropout)	(None, 8, 8, 64)	0
conv2d_16 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_18 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_17 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_19 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_10 (Dropout)	(None, 4, 4, 128)	0
flatten_2 (Flatten)	(None, 2048)	0
dense_4 (Dense)	(None, 128)	262272
batch_normalization_20 (Batch Normalization)	(None, 128)	512
dropout_11 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 10)	1290




Second Attempt: Transfer Learning




- Introduction: Taking a pre-trained neural network and adapting it to a new task
- Popular models: GoogLeNet network, ImageNet
- Results:

- val_output_accuracy: 0.2540

Test accuracy: 0.7008000016212463





Third Attempt: Learning Schedule and Regularization




- Learning Schedule: Also known as learning rate schedule or learning rate decay, involves adjusting the learning rate during training

```
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate,  
    decay_steps=train_images.shape[0] // batch_size,  
    decay_rate=0.95,  
    staircase=True)
```

- Regularization: Adding additional constraints to the training process to prevent overfitting


Method: Dropout

```
layers.Conv2D(64, (3, 3), activation='relu', padding='same',  
    kernel_regularizer=l2_reg),  
layers.BatchNormalization(),  
layers.MaxPooling2D((2, 2)),  
layers.Dropout(0.5),
```








Final Model Structure



Final Model Accuracy



Accuracy: 0.8715999722480774

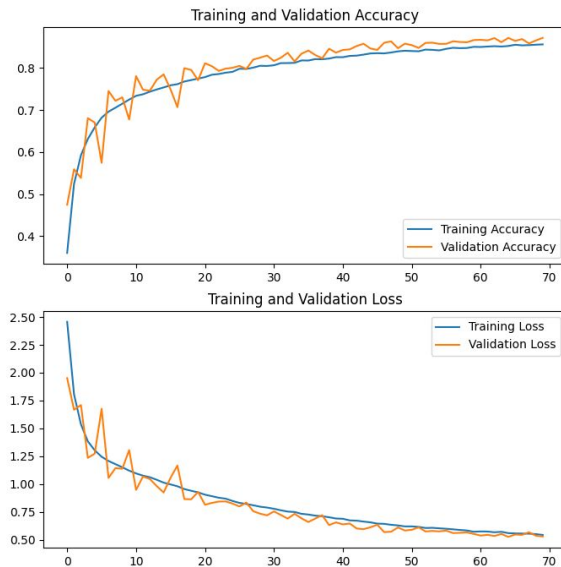
Way to Find the Maximum Epochs Needed

1. Pseudocode:

```
Epoch = 1000
Run model
Count == 0
While {
  Track every accuracy outputted by Model
  If ( Accuracy_{i} - Accuracy_{i-1} <= N )
    Count += 1
  If (count == 3)
    Terminate model
    Break}
```


2. Example:

```
Epoch 11/50
781/781 [=====] - 532s 681ms/step - loss: 1.1138 - accuracy: 0.7327 - val_loss: 1.0703 - val_accuracy: 0.7528
Epoch 12/50
781/781 [=====] - 511s 655ms/step - loss: 1.0839 - accuracy: 0.7361 - val_loss: 0.9837 - val_accuracy: 0.7686
```

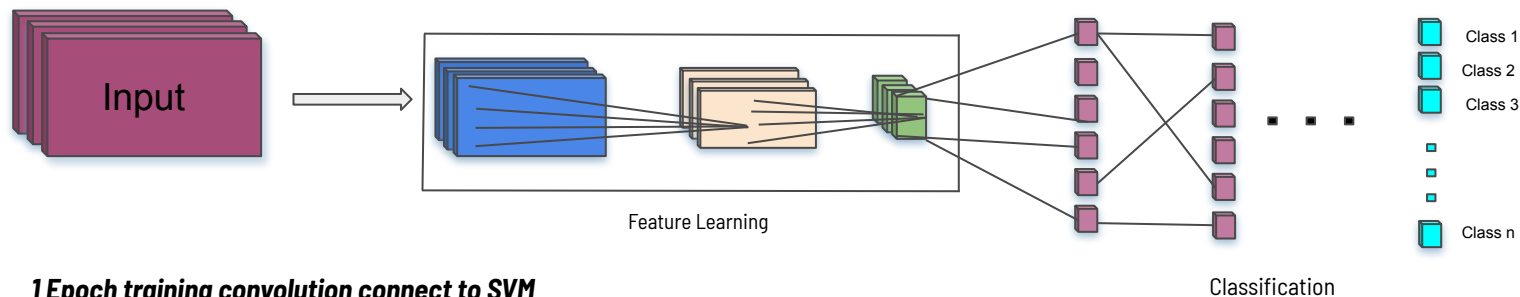


Convolutional SVM

We used CNN as a socket for SVM, and the flowchart would be shown as below:



CNN Flowchart:




1 Epoch training convolution connect to SVM

CNN accuracy:0.4151

Convolutional SVM accuracy:0.5797

Conclusion



Performance (SVM vs. CNN)

Why ? (SVM suited for simpler datasets with fewer dimensions, CNN works better for image recognition task)





Thanks!

Do you have any questions?

